

# M.Sc. IN HIGH-PERFORMANCE COMPUTING

## MAP55672 - CASE STUDIES IN HPC

### COMMUNICATION-AVOIDING QR FACTORIZATION

Kirk M. Soodhalter (ksoodha@maths.tcd.ie)  
School of Mathematics, TCD

#### RULES

---

Please create a git repository with your solutions to this assignment, and submit the link to me by email. One folder in your git repository should contain a Matlab Live file or Jupyter notebook file with the part of the assignment you are asked to do in an interpreted language. The second folder should contain all your C code implementing CAQR the second part of the assignment with a makefile the runs everything. by **Friday, 21. February, 2025.**

#### QUESTION

---

In this assignment, you are being asked to build a simple communication-avoiding TSQR-factorisation of a tall, narrow matrix in both Matlab/Python and then in C programming language. Some parameters:

- We restrict our focus to the case that the tall, narrow matrix is distributed onto four processors, like in the example from the lectures;
  - For the MATLAB part of the assignment, you will not do any parallel programming; you are simply meant to work out the linear algebra of communication-avoiding TSQR-factorisation; you should take advantage of these languages abilities to directly address sub-blocks of a matrix using *sub-indices*. For example, in MATLAB,  $W(i : j, :)$  returns rows  $i$  to  $j$  of the matrix  $W$ . You may use built-in methods to perform local QR-factorisations of the blocks. In MATLAB, this would be  $[Q, R] = \text{qr}(W\_bl)$
  - For the C programming part of the assignment, you should familiarize yourself with the QR factorisation routines contained in LAPACK. You will use the Householder-based QR LAPACK routine to perform local, on-processor QR factorisations.
1. [30 points] In MATLAB or Python, implement a version of the TSQR that divides an input matrix up into four blocks of rows (using row sub-indexing) and computes the QR-factorisation in the way shown in the lectures on communication-avoiding factorisations.

2. [50 points] In C programming language, implement a method called `TSQR()` which executes just the communication-avoiding QR-factorisation of a tall-narrow matrix using the technique described in the lecture videos/included technical report. It should use LAPACK's QR factorisation routine for local QR. Rows of the matrix should be evenly distributed onto four compute nodes, as described in the lectures. Set up a (relatively) small test problem to demonstrate that your code works (i.e., that the factorisation is correct).
3. [20 points] Generate a series of random  $m \times n$  test matrices for different values of  $m$  and  $n$ . How does your code scale with respect to each of these dimensions? Please show some scaling plots to answer this question.